

Education

University of Toronto

Master of Science in Applied Computing, Artificial Intelligence Concentration

Sept '23 – Dec '24 (expected)

- Vector Scholarship in Artificial Intelligence, worth \$17 500 CAD.

University of Waterloo

Bachelor of Computer Science, Honours, Co-operative Program, AI Specialization

Sept '18 – Apr '23

- 93.50 major average, with courses in computer vision, autonomous vehicles, and machine learning.
-

Peer-Reviewed Publications

Sim-to-Real Domain Adaptation for Lane Detection and Classification in Autonomous Driving

June 2022

- Co-authored a peer-reviewed [research paper](#) on sim2real lane detection using the [CARLA](#) simulator, integrating current GAN research with state-of-the-art lane detection models.
- Presented the paper at a poster board session during the [Intelligent Vehicles 2022 symposium](#).

Application of Machine Learning for Drone Classification Using Radars

April 2021

- Authored a research paper and presented it at the [SPIE 2021 Conference](#), with a follow-up [peer-reviewed paper](#).
 - Used a convolutional neural network to classify 5 different commercial drones based on a Fourier transform of their noisy RADAR return signal.
 - Created a detailed physics-based simulation of RADAR returns bouncing off of a drone in flight, and used the simulated data to validate classification methods.
-

Relevant Work Experience

Cruise

SAN FRANCISCO, CA

Model Deployment Platform Engineer

Sept '22 – Dec '22

- Designed and developed debugging tools to assist in model deployment, reducing deployment time by days in some cases.
- Built a tool to automatically assess numerical divergence between the original PyTorch models and optimized [TensorRT](#) models, allowing deployment engineers to quickly identify optimization errors.
- Built a linter to map [ONNX](#) nodes to lines in PyTorch source code, saving deployment engineers hours of manual debugging every deployment.
- Created a tool to manually add outputs to exported ONNX graphs, allowing engineers to inspect intermediate activations while debugging models.

NVIDIA

SANTA CLARA, CA (REMOTE)

Deep Learning Researcher for Autonomous Vehicles

Jan '22 – Apr '22

- Designed and iterated on multiple experiments for a LiDAR object detection neural network, improving cyclist and pedestrian detections by 43% and 15%, respectively.
- Implemented sparse tensor object detectors using [Minkowski Engine](#), outperforming the baseline model while using 70% less memory.
- Presented a 30-minute research overview to the organization, communicating 4 months worth of findings and insights to internal stakeholders and leadership.
- Integrated confidence predictions into a LiDAR object detection auto-labeling pipeline, allowing human annotators to focus efforts on anomalous and challenging data instances.

DarwinAI

WATERLOO, ON (REMOTE)

Machine Learning Developer

May '21 – Aug '21

- Built and tested defect detection deep learning solutions for clients in the manufacturing industry, working in cross-functional teams and presenting results to stakeholders weekly.
- Trained XGBoost and SVR systems to model the relationship between environmental conditions and yield for an agriculture client, giving the client insight into the impact of environmental conditions on their yield.
- Implemented the core functionality of Dataset Distillation using the [autograd](#) package, to pursue research in low-data machine learning contexts.
- Implemented VAE, VQ-VAE, and VQ-VAE-2 from scratch in PyTorch, demonstrating the effectiveness of autoencoders as anomaly detectors on client datasets.

Untether AI

TORONTO, ON (REMOTE)

Software Developer

Sept '20 – Dec '20

- Contributed to a customer-facing Python library to quantize and optimize TensorFlow computation graphs, towards compiling models for custom hardware.
- Designed and implemented Non-Max Suppression for quantized values using only integer operations, eliminating a CPU bottleneck and allowing object detection pipelines to run 10x faster on-chip.
- Experimented with different quantization schemes to improve object detection performance of an SSD-ResNet-34 by 5%.
- Implemented a lookup table class to represent arbitrary non-linear functions in a quantized space, allowing the Python library to support dozens of additional model architectures.

Huawei

MARKHAM, ON

LiDAR Perception Researcher

Jan '20 – Apr '20

- Implemented key modules from 12 different research papers in PyTorch, summarizing state-of-the-art techniques and enabling further research.
- Analysed the SemanticKITTI dataset to produce optimal class loss weights, boosting mIoU by 2%.
- Wrote a data loader to spatially align sequential LiDAR scans for temporal pipelines, based on IMU data, enabling the research of multi-scan LiDAR segmentation pipelines.

Please refer to my [Linkedin profile](#) for a complete list of work experiences.

Skills

Languages: Python, C++, R, C, LaTeX, Java, JavaScript

Frameworks: PyTorch, TensorFlow, NumPy, Pandas, OpenCV, ROS, Gradio

Tools: Git, Docker, Conda, Bazel, CARLA, GCP, Linux, TensorRT